

COMPLEX-SYSTEMS-1

October 18, 2024

1 Zipf's law

1.1 Author - Damian Pietroni

1.2 Introduction

Zipf's law is a phenomenon that states that, given a set of occurrences ordered in decreasing order, the value of the n -th element is inversely proportional to its frequency of occurrence.

The mathematical representation of Zipf's law can be written as:

$$f(n) \propto 1/n$$

where:

- $f(n)$ is the frequency of the n -th element,
- n is the rank of the element.

1.3 Purpose of this lab

This laboratory aims to explore and gain a deeper understanding of Zipf's Law and the phenomena it describes in large text body. By analyzing the frequency distribution of words in extensive texts, we will observe how Zipf's Law aligns with a frequency distribution across text

1.4 Code overview

Code for this lab was written both in Python and C++, Algorithmic part was done with C++, but plots were done with Python

```
[ ]: pip install -r requirements.txt
```

In order to download all necessary python libraries run command above (first remember to create virtual environment)

1.5 Import libraries

```
[187]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from pathlib import Path
```

```
[188]: # load all csv
path = Path('data')
files = path.glob('*.csv')
```

```
[189]: # Load the data
data_frames = {}
for file in files:
    file_name = file.name.replace('.csv', '')
    print(file_name)
    data = pd.read_csv(file)
    data_frames[file_name] = data
```

```
bible
divine_comedy
faust_de
LLM_de
plato
_1984
```

1.5.1 Why do I need columns 3 and 4?

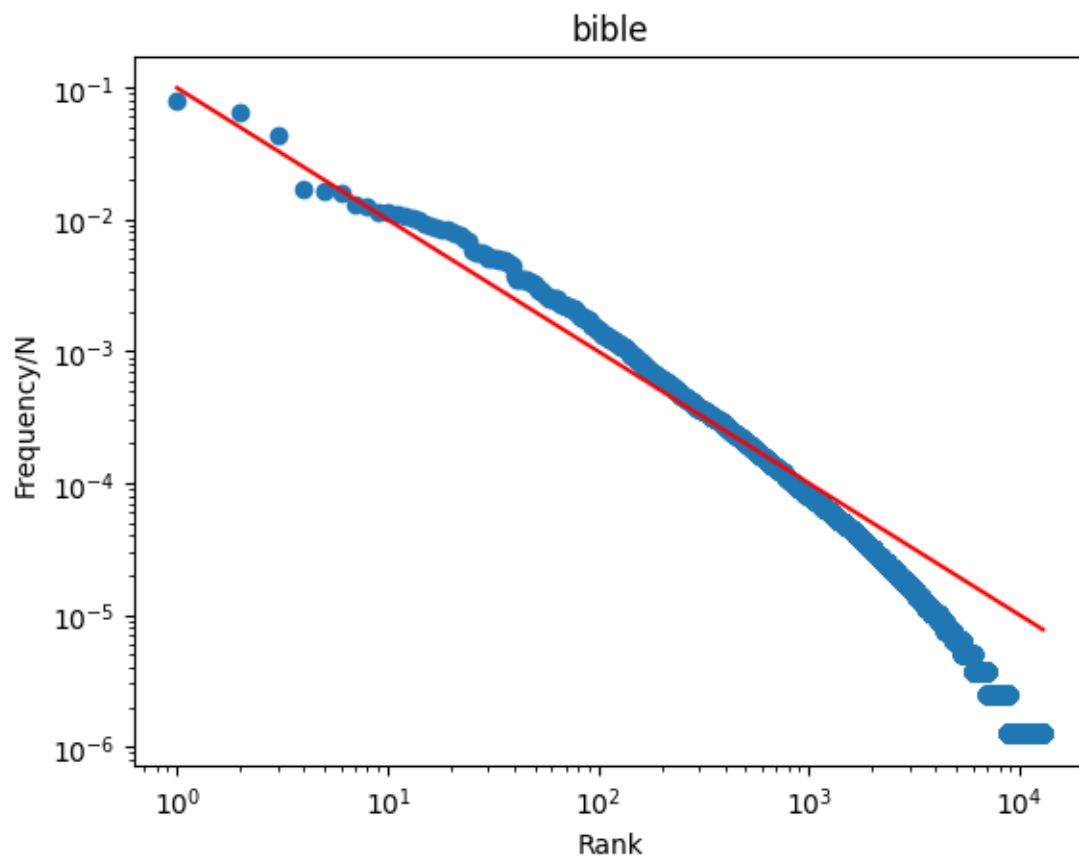
- Column 3 (Rank) is needed to calculate Theoretical ZipfLaw
 - Column 4 (Frequency) is needed to calculate real frequency of a word
-

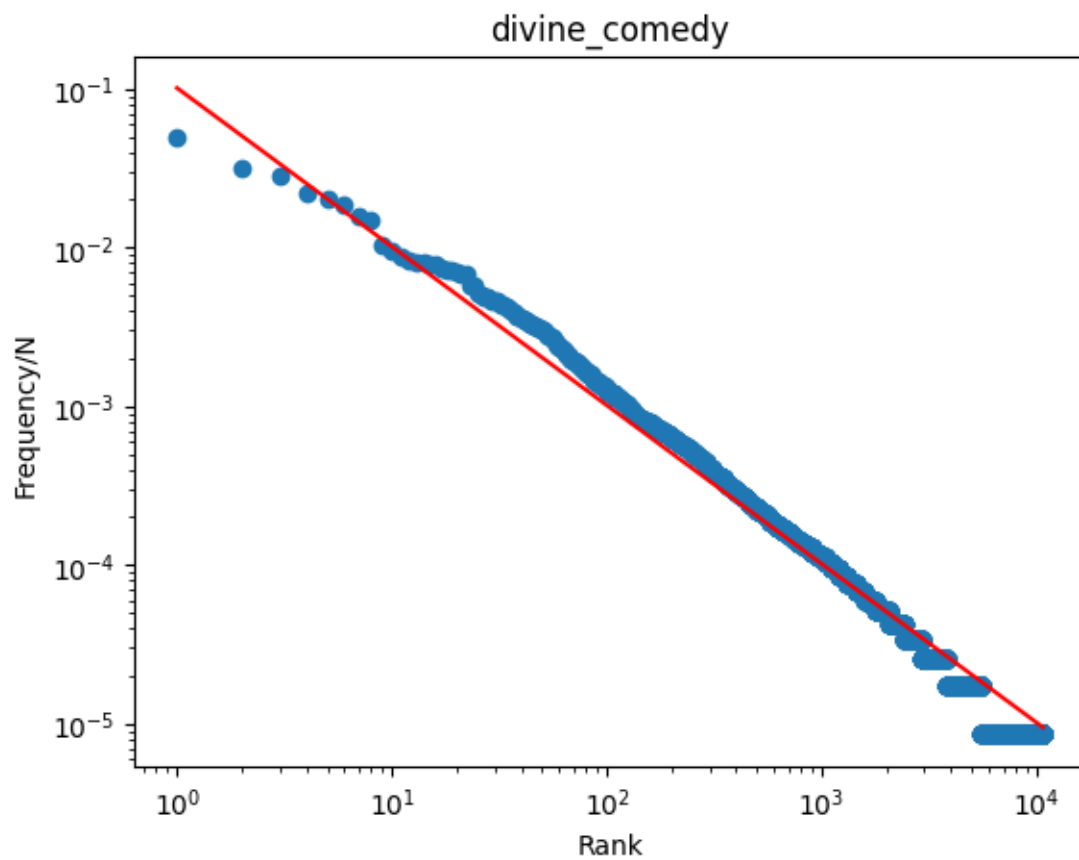
1.6 All books with log to log comparison

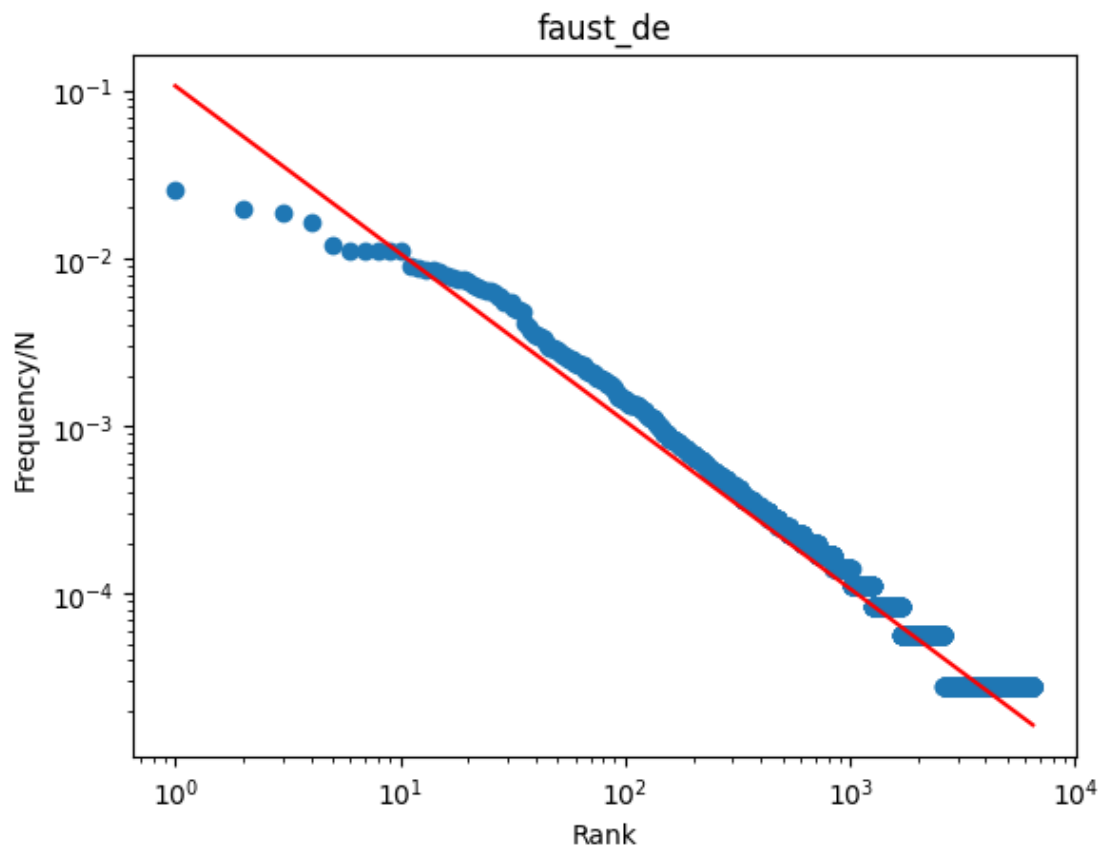
```
[190]: fig = plt.figure(figsize=(4, 1))
for book_name, book in zip(data_frames.keys(), data_frames.values()):
    plt.figure()
    truncated_data = data_frames[book_name]
    plt.scatter(truncated_data['Rank'], truncated_data['Frequency/N'],
    ↪marker='o', label=book)
    plt.plot(truncated_data['Rank'], truncated_data['Zipf'], label='Zipf',
    ↪color='red')
    plt.xscale('log')
    plt.yscale('log')
    plt.xlabel('Rank')
```

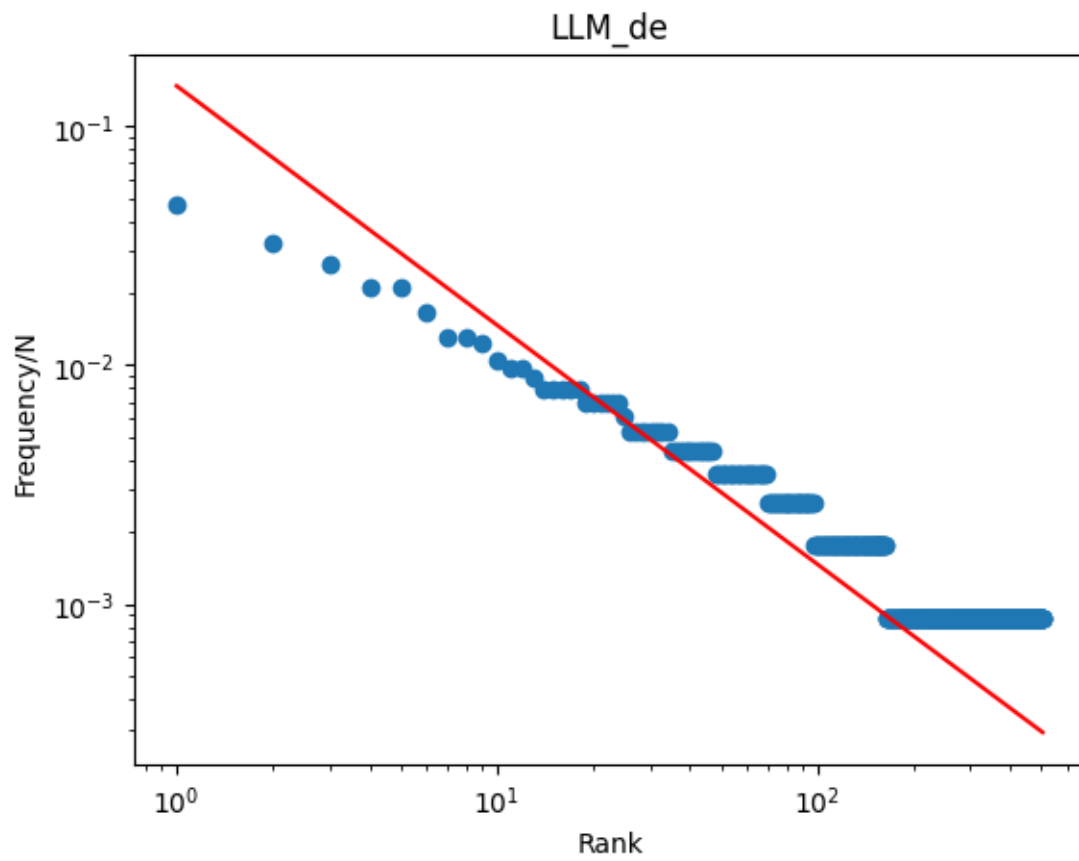
```
plt.ylabel('Frequency/N')  
plt.title(book_name)  
plt.show()
```

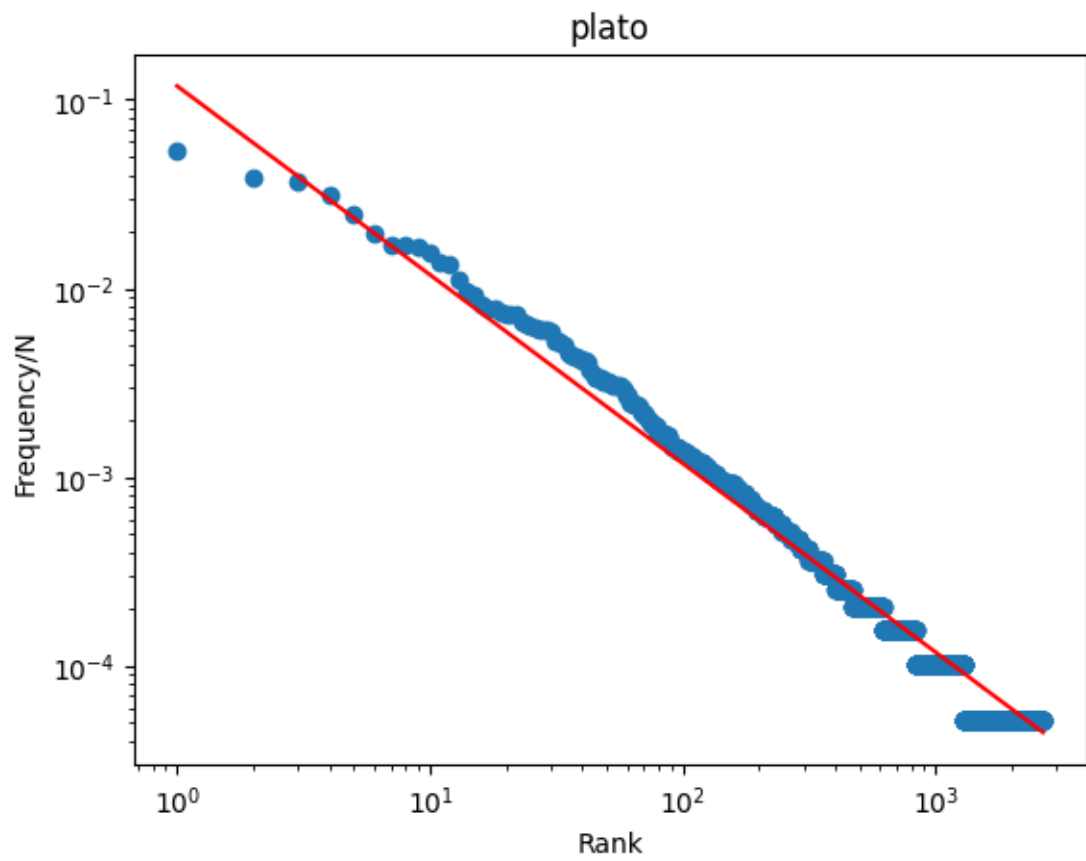
<Figure size 400x100 with 0 Axes>

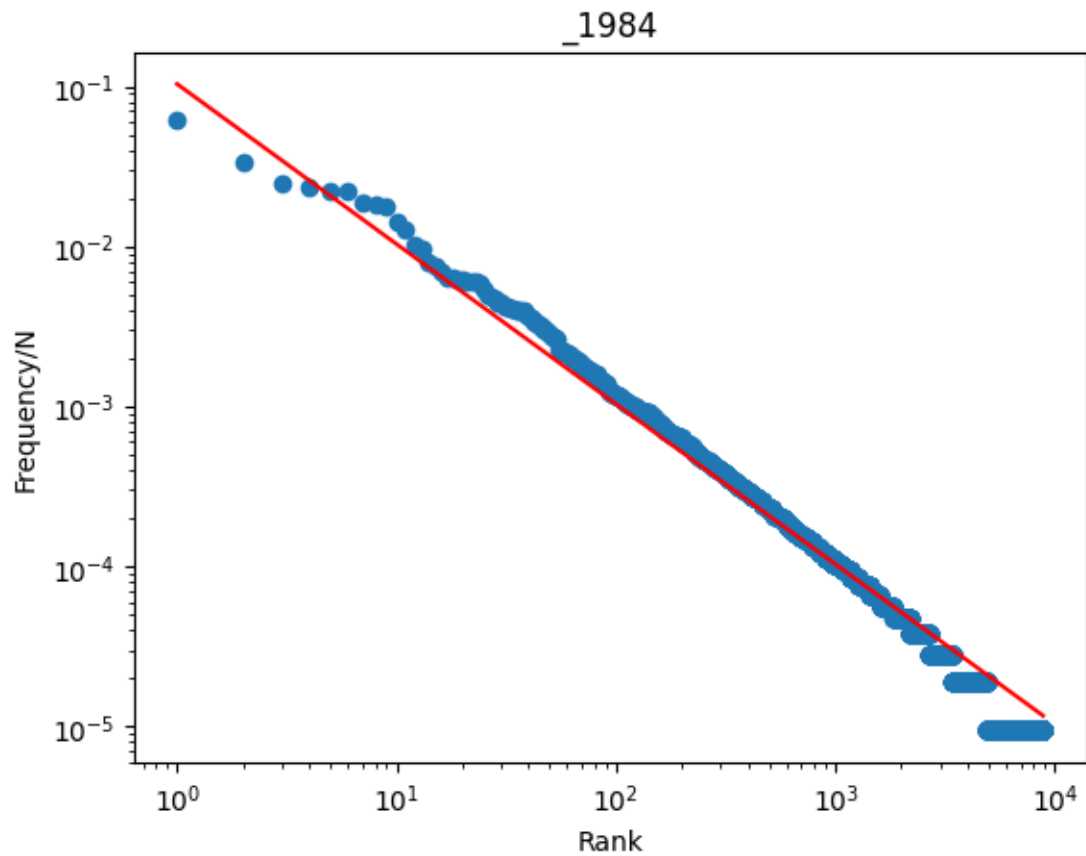








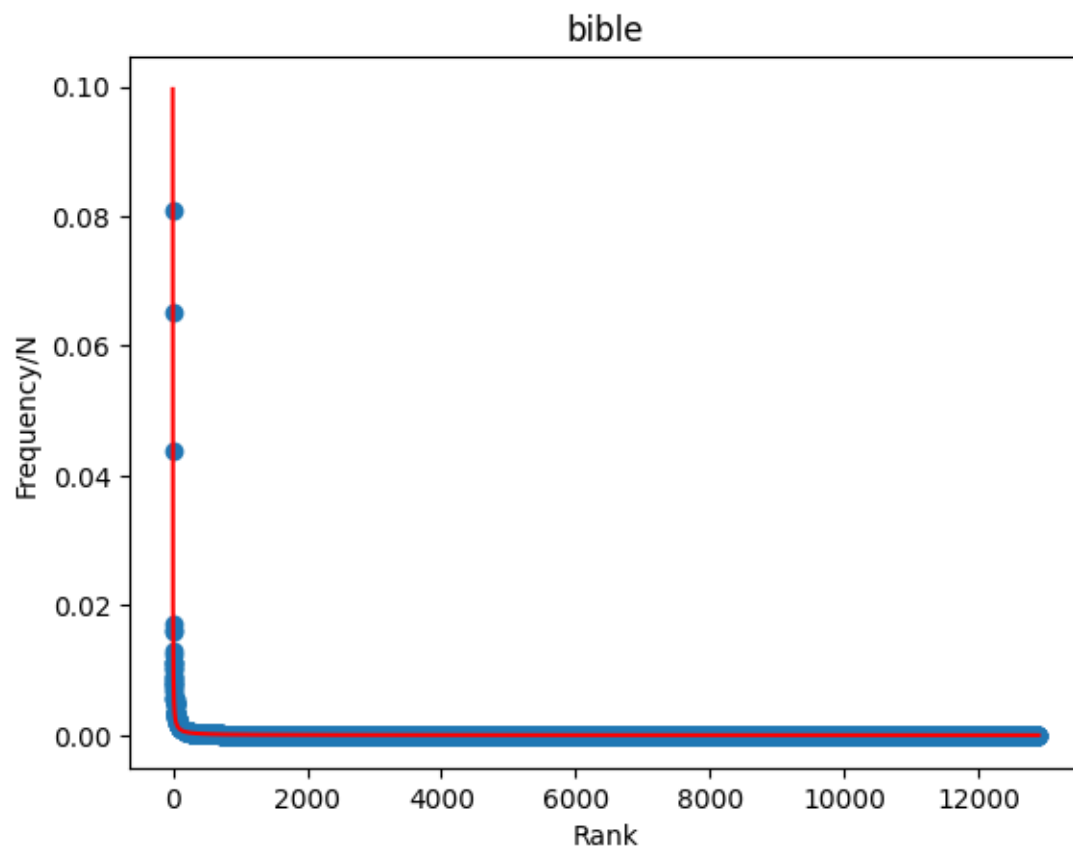


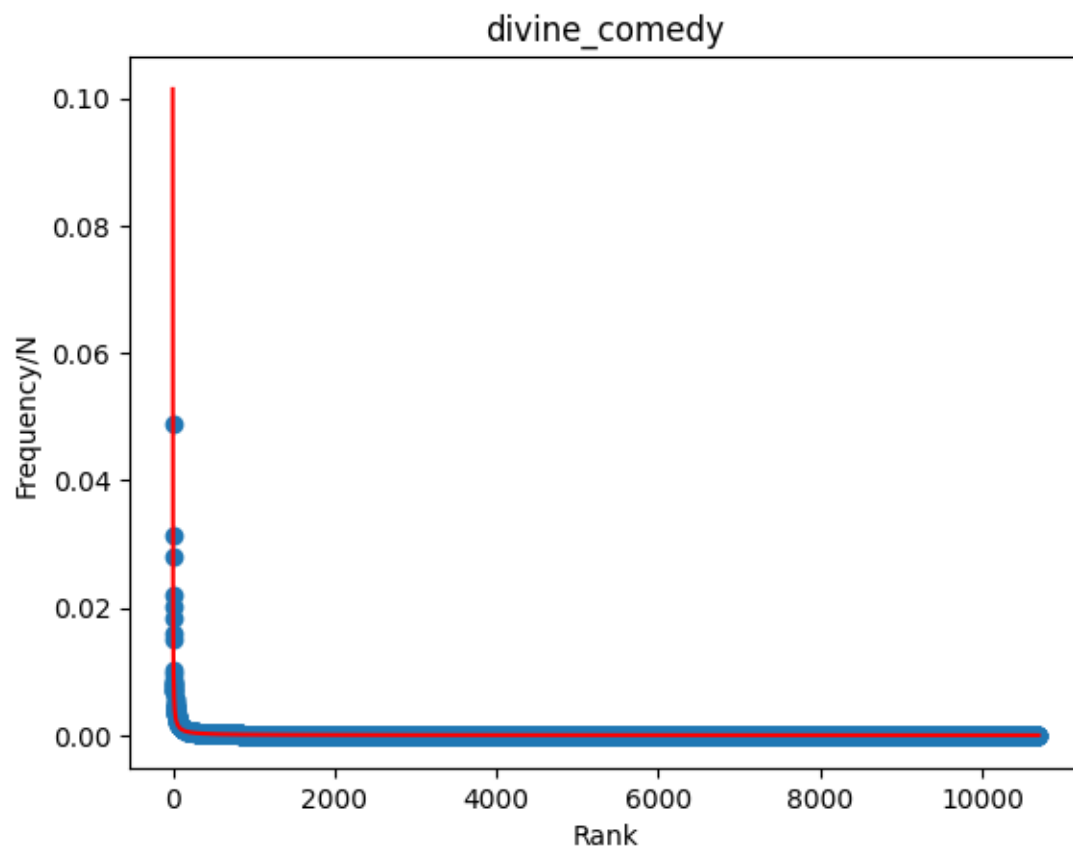


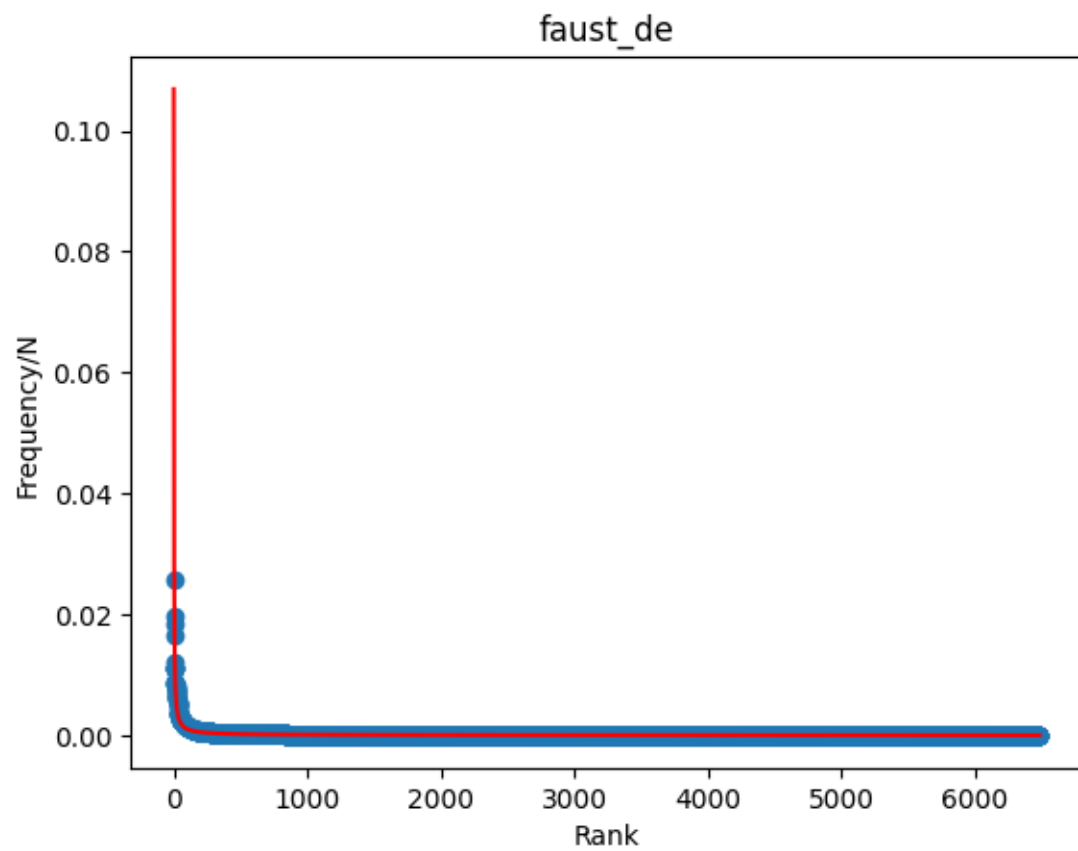
1.7 All books with linear comparison

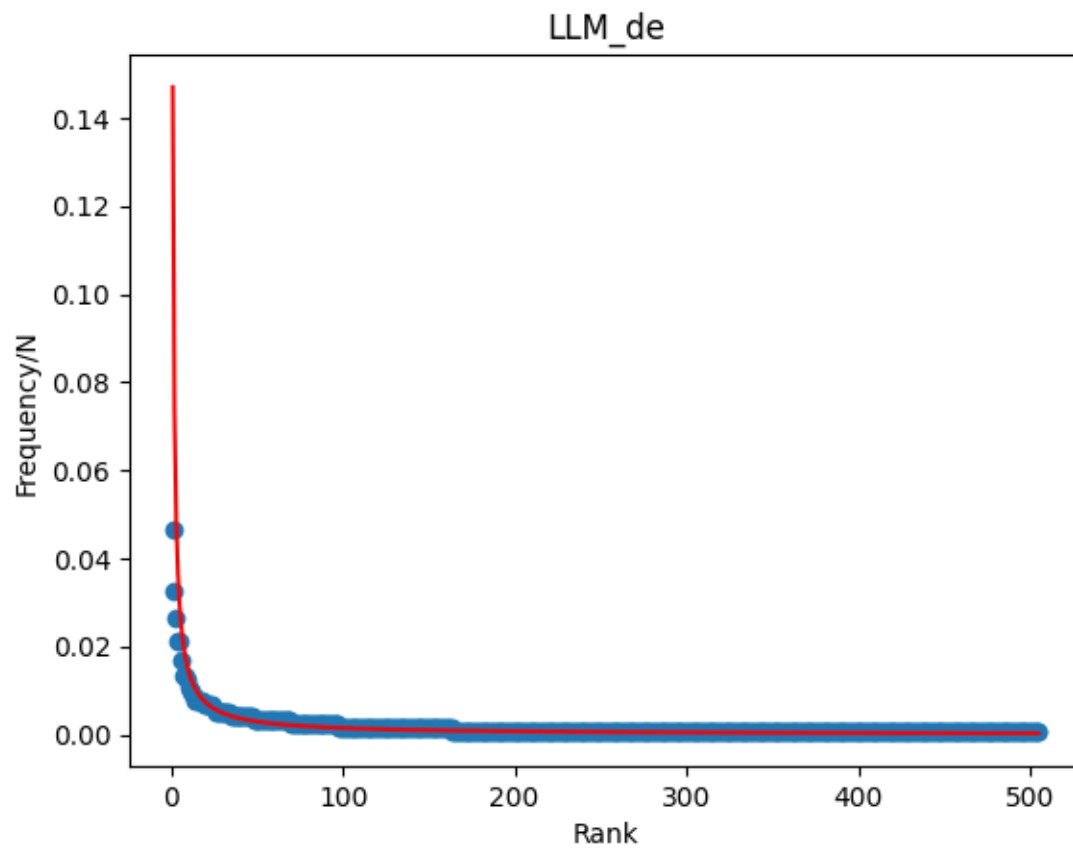
```
[191]: fig = plt.figure(figsize=(4, 1))
for book_name, book in zip(data_frames.keys(), data_frames.values()):
    plt.figure()
    truncated_data = data_frames[book_name]
    plt.scatter(truncated_data['Rank'], truncated_data['Frequency/N'],
    ↪marker='o', label=book)
    plt.plot(truncated_data['Rank'], truncated_data['Zipf'], label='Zipf',
    ↪color='red')
    plt.xlabel('Rank')
    plt.ylabel('Frequency/N')
    plt.title(book_name)
    plt.show()
```

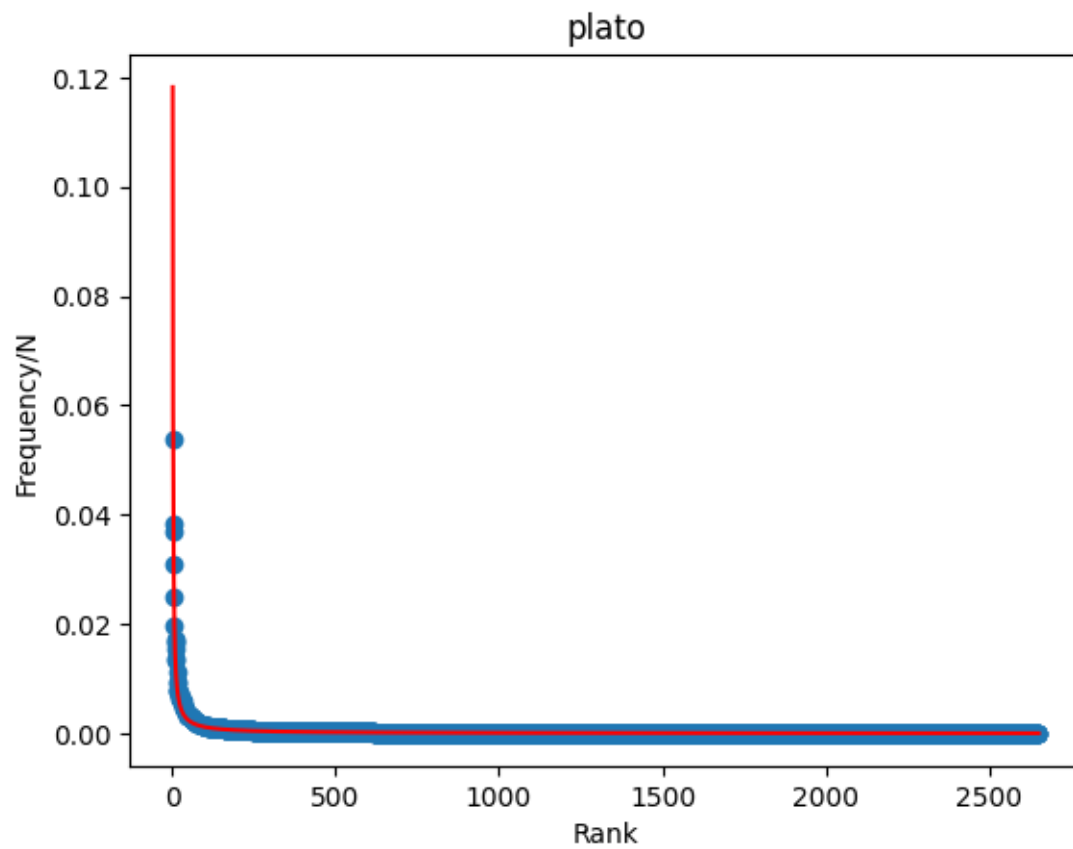
<Figure size 400x100 with 0 Axes>

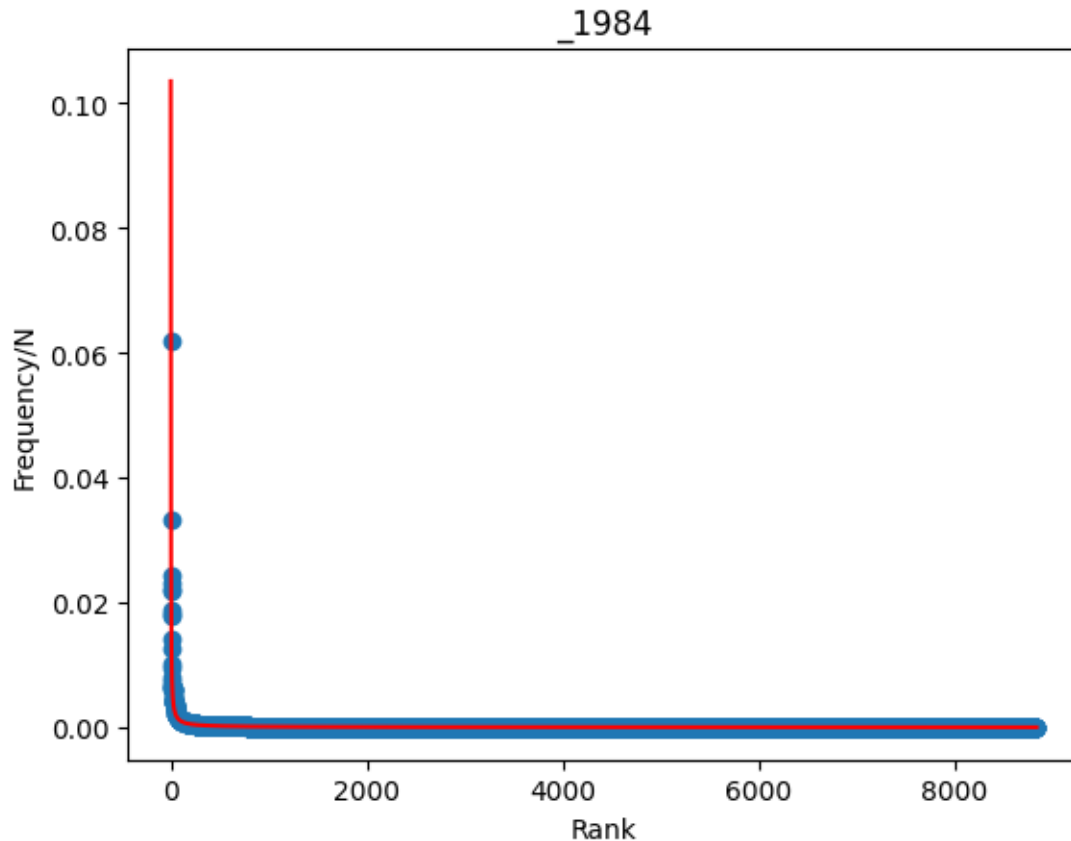












When comparing data on a linear scale, subtle changes are often difficult to observe, which become more apparent when viewed on a log-log scale. This is because **ranks increase exponentially**, typically by a factor of $10e4$. Since Zipf's Law involves the inverse of ranks, it also spans a similarly large range. In a log-log scale, this wide range is compressed and more manageable, making patterns and relationships easier to detect, whereas on a linear scale, the data can be too spread out to reveal these insights effectively.

1.8 Extra 1 & Extra 2

```
[192]: path = Path('data_fitted')
       files = path.glob('*.csv')
```

```
[193]: data_frames = {}
       for file in files:
           file_name = file.name.replace('.csv', '')
           print(file_name)
           data = pd.read_csv(file)
```

```
data_frames[file_name] = data
print(data.head(2))
```

bibleFitted

	Token	Frequency	Rank	Frequency/N	Zipf	FittedZipf	A	B
0	the	64193	1	0.080766	0.099588	0.082163	0.9	0.6
1	and	51763	2	0.065127	0.049794	0.053077	0.9	0.6

divine_comedyFitted

	Token	Frequency	Rank	Frequency/N	Zipf	FittedZipf	A	B
0	the	5721	1	0.048968	0.101467	0.050733	1	1
1	d	3655	2	0.031284	0.050733	0.033822	1	1

faust_deFitted

	Token	Frequency	Rank	Frequency/N	Zipf	FittedZipf	A	B
0	und	917	1	0.025841	0.106898	0.022068	0.9	6.4
1	ich	700	2	0.019726	0.053449	0.019689	0.9	6.4

LLM_deFitted

	Token	Frequency	Rank	Frequency/N	Zipf	FittedZipf	A	B
0	die	53	1	0.046491	0.146999	0.044545	1	2.3
1	der	37	2	0.032456	0.073499	0.034186	1	2.3

platoFitted

	Token	Frequency	Rank	Frequency/N	Zipf	FittedZipf	A	B
0	the	1047	1	0.053994	0.118197	0.052859	0.9	2.1
1	of	744	2	0.038368	0.059098	0.041100	0.9	2.1

_1984Fitted

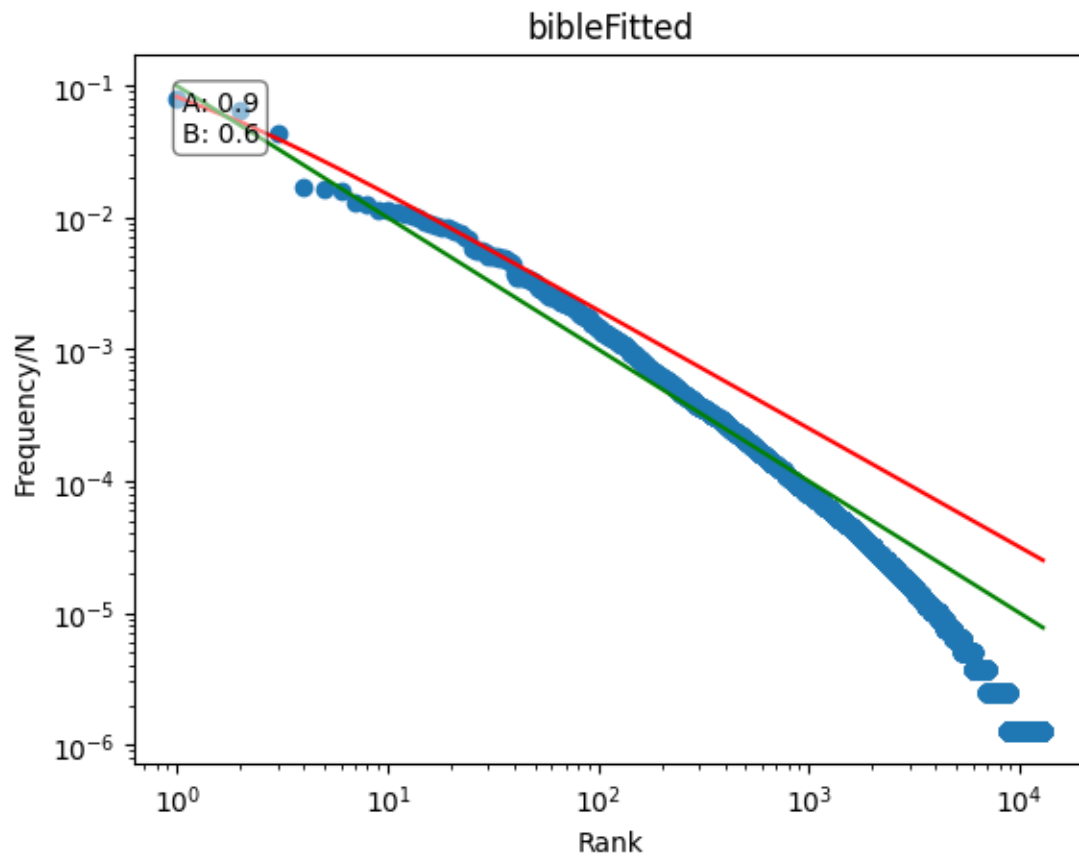
	Token	Frequency	Rank	Frequency/N	Zipf	FittedZipf	A	B
0	the	6527	1	0.061826	0.103473	0.060867	1	0.7
1	of	3500	2	0.033153	0.051737	0.038324	1	0.7

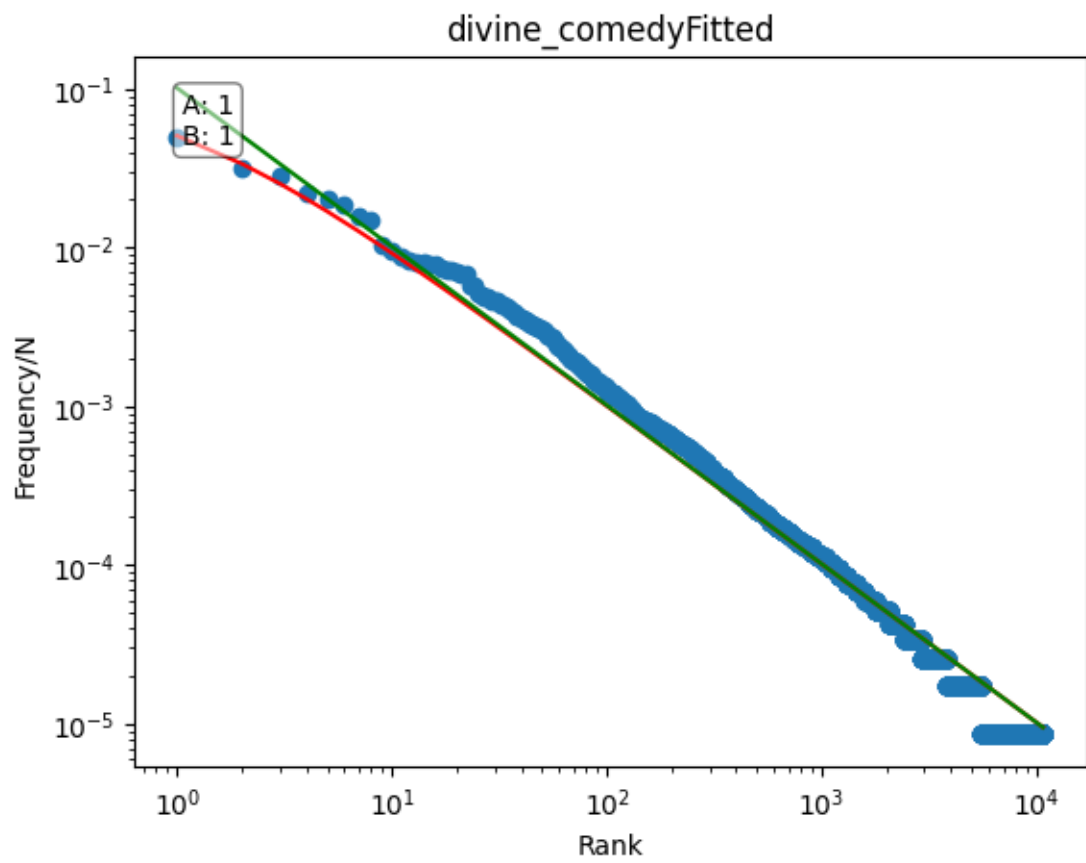
```
[194]: fig = plt.figure(figsize=(4, 1))
for book_name, book in zip(data_frames.keys(), data_frames.values()):
    truncated_data = data_frames[book_name]
    plt.figure()
    plt.scatter(truncated_data['Rank'], truncated_data['Frequency/N'],
    ↪marker='o', label=book)
    plt.plot(truncated_data['Rank'], truncated_data['FittedZipf'],
    ↪label='FittedZipf', color='red')
    plt.plot(truncated_data['Rank'], truncated_data['Zipf'], label='Zipf',
    ↪color='green')
    a_value = truncated_data['A'].values[0]
    b_value = truncated_data['B'].values[0]
    plt.text(0.05, 0.95, f"A: {a_value}\nB: {b_value}", transform=plt.gca().
    ↪transAxes, fontsize=10,
              verticalalignment='top', horizontalalignment='left',
    ↪bbox=dict(boxstyle='round', facecolor='white', alpha=0.5))

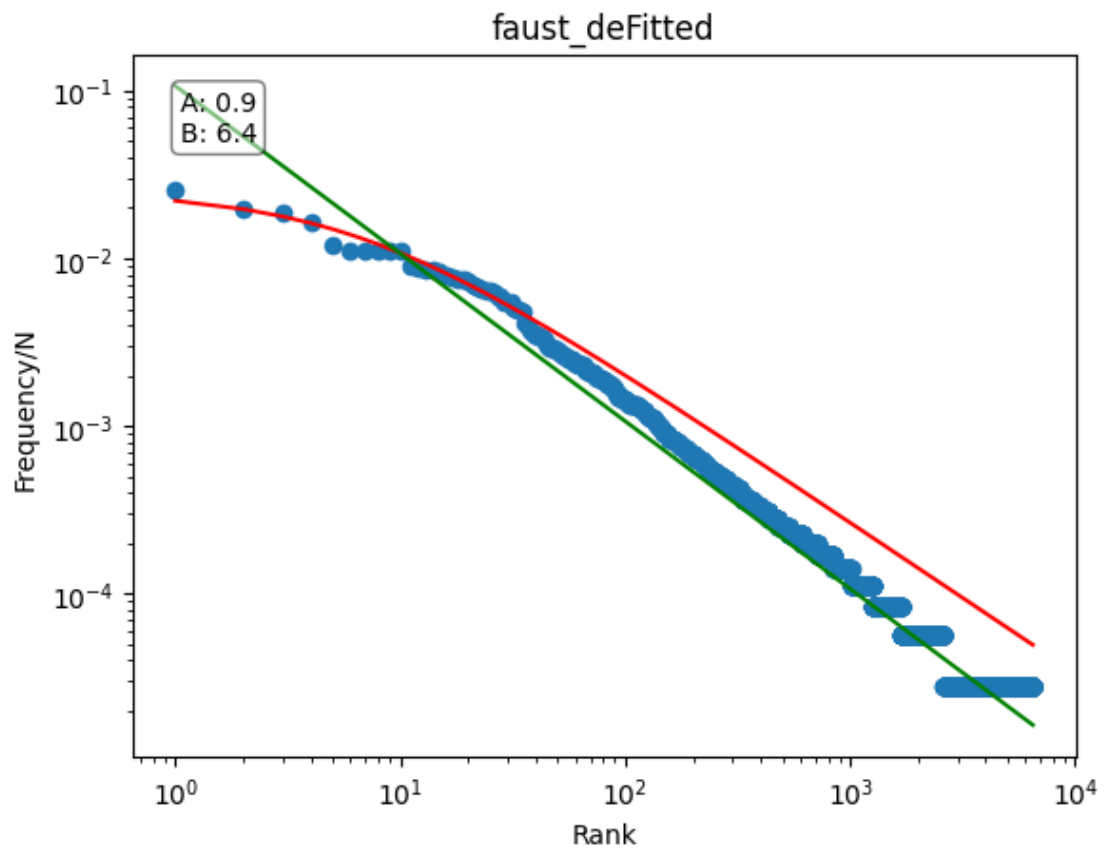
    plt.xscale('log')
    plt.yscale('log')
```

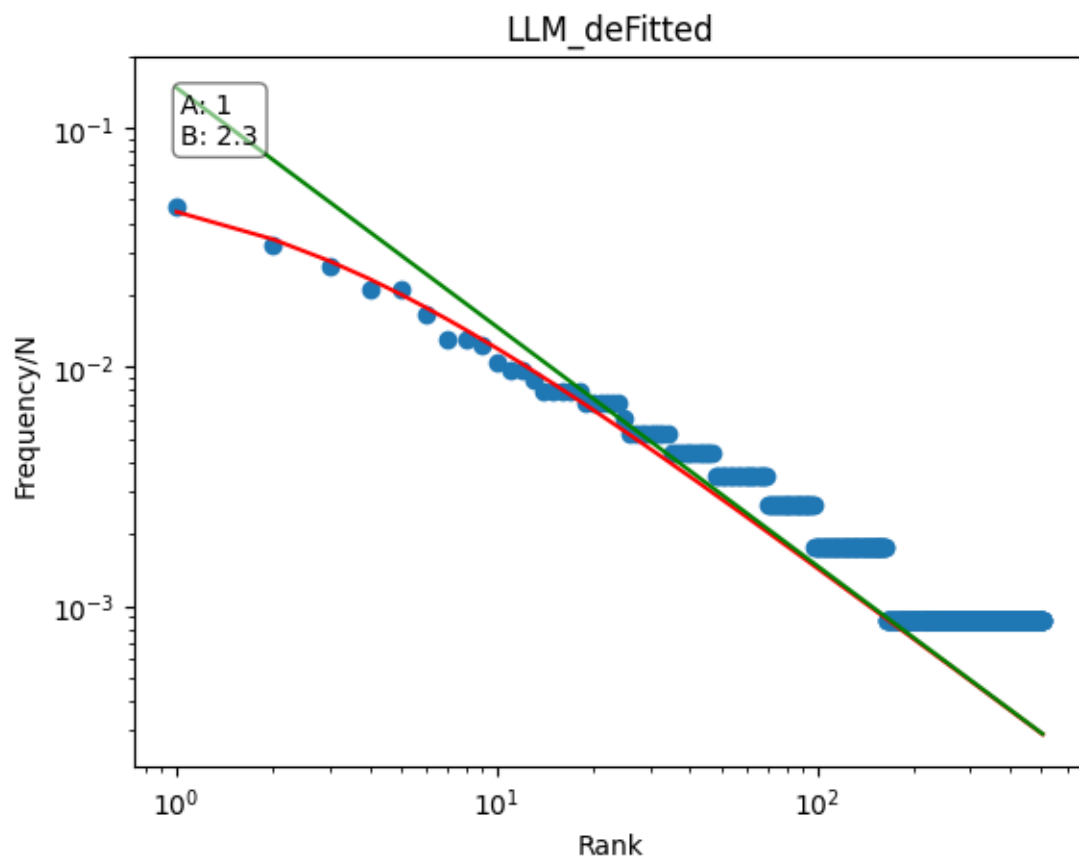
```
plt.xlabel('Rank')
plt.ylabel('Frequency/N')
plt.title(book_name)
plt.show()
```

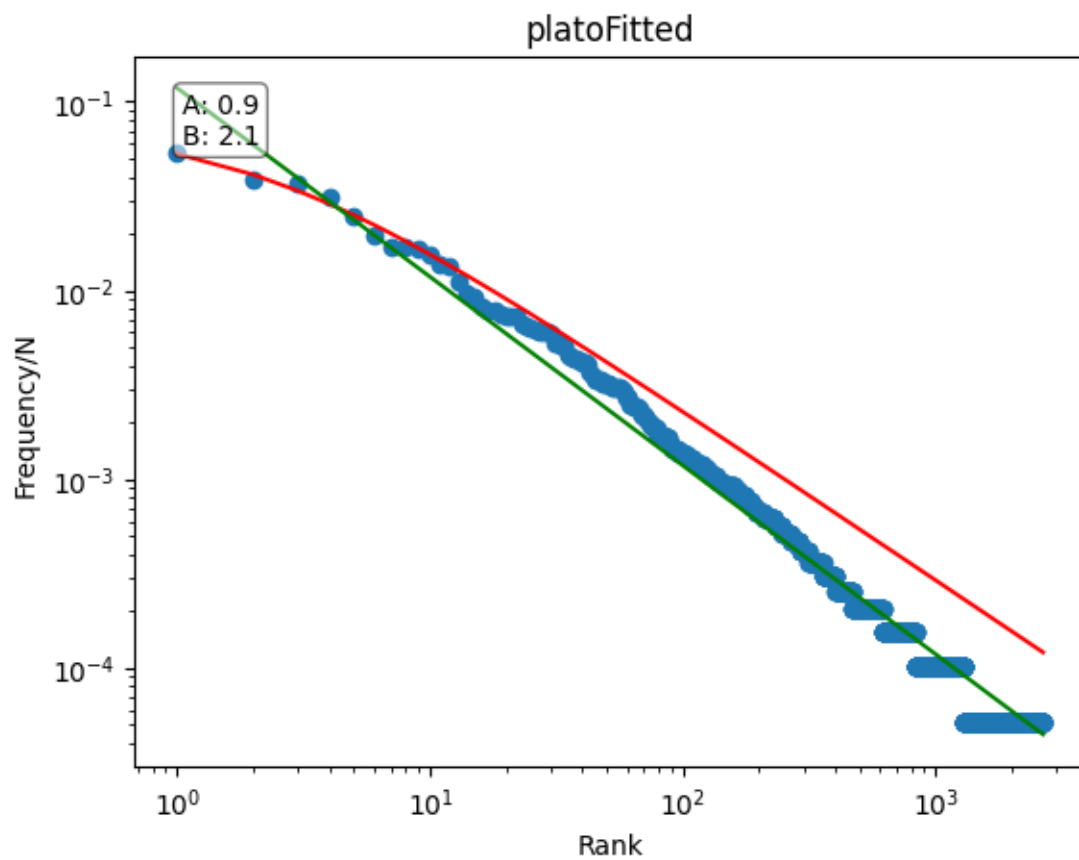
<Figure size 400x100 with 0 Axes>

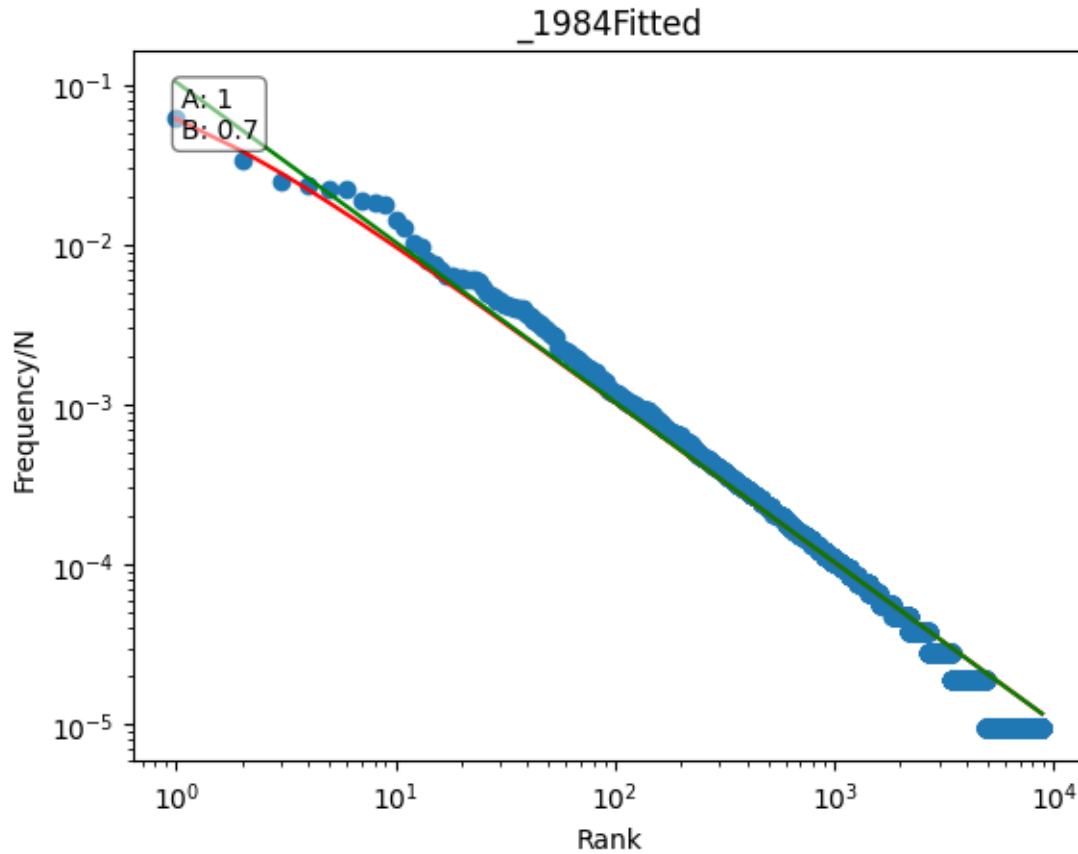












As we can observe, the constants a and b vary across different texts. However, I believe that relying solely on these constants would not be sufficient to determine the language used in a book. It might be just a suggestion.-

The output of language models (LLMs) can be analyzed through the lens of the Zipf-Mandelbrot law. The values obtained can be fitted using linear regression, indicating that their distribution resembles that of Zipf's law. While the fit may not appear as precise as in other plots, this could be attributed to the limited sample size of the words.

Code: <https://github.com/neuropython/ComplexSystems/blob/master/zipfLawAlgorithms.cpp>